

Polyglot data system

applied to



Francesco Foresi | 508971
Pier Vincenzo De Iellis | 512521



MB

Data Management Problems

- ❑1 Various types of data model
- ❑2 No scalable architecture
- ❑3 No common view of data resources


Data Management Solution



Multiple
Data
stores



Query
System



Heterogeneous
Storage
Engine



“One size does not fit all”

–State of Art



Federated vs Polystore

- Unique Query Language for all DBs
- Simple implementation
- Individual storage engines are independent

Federated

VS

- Native Query Language for DBs
- Hard implementation
- Storage engines are managed together as an integrated set

Polystore


BigDAWG vs Polypheny

- Different way of organizing data and making queries
- Different data migration
- Different Architecture

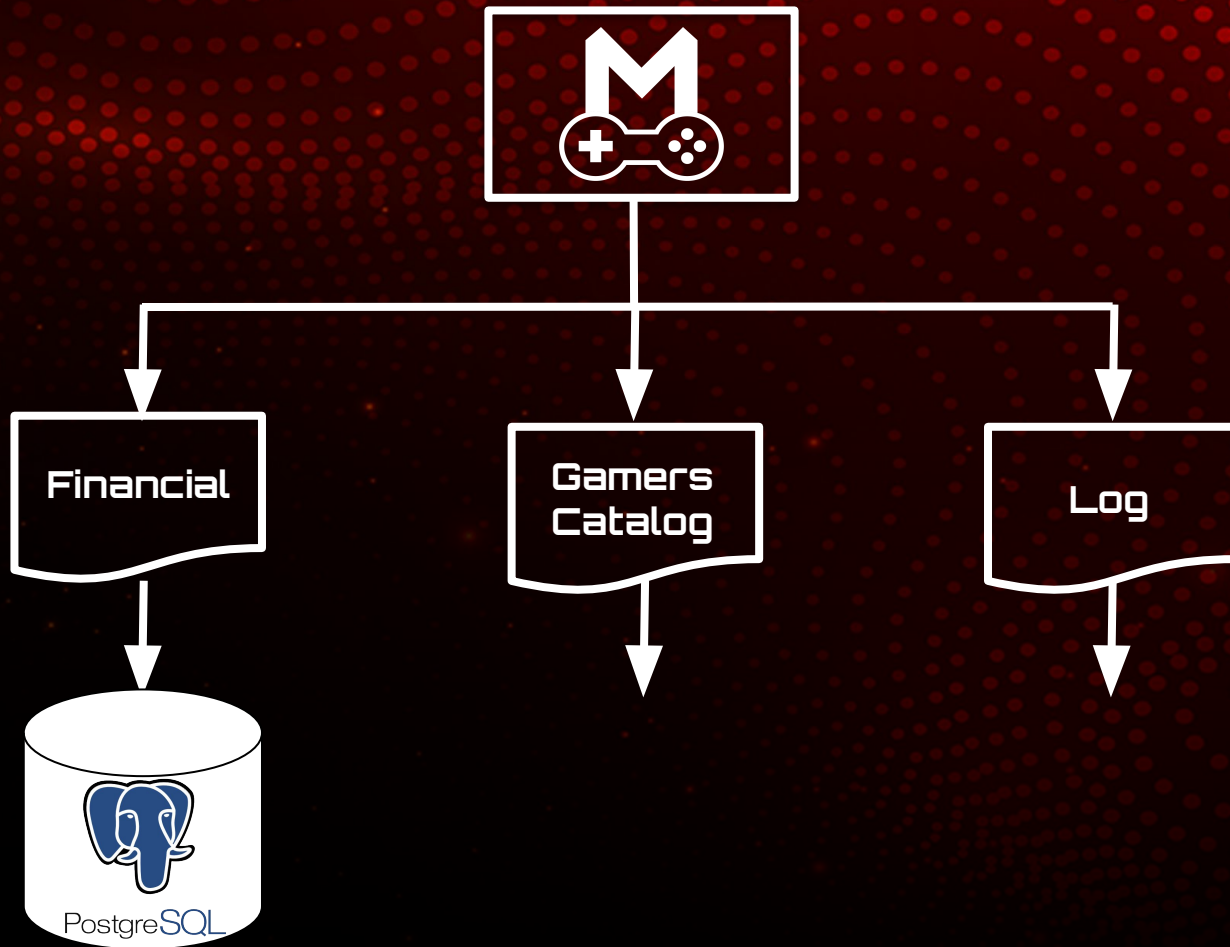
Our choice: BigDAWG

- **BigDAWG** is based on 3 fundamental concepts:
 - **Islands** (data model and programming model)
 - **Shim** (model translate)
 - **Cast** (data migration)

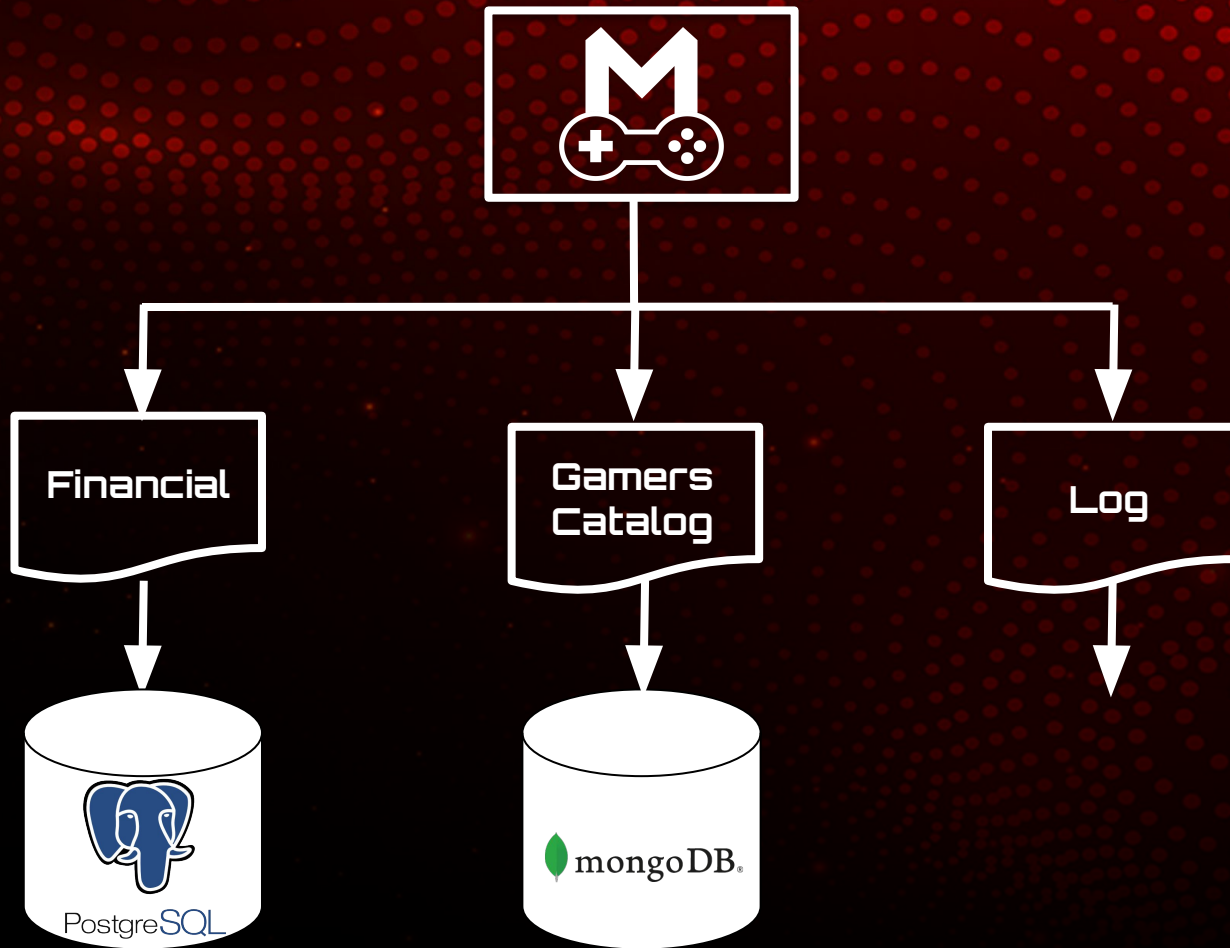
M8 Polystore



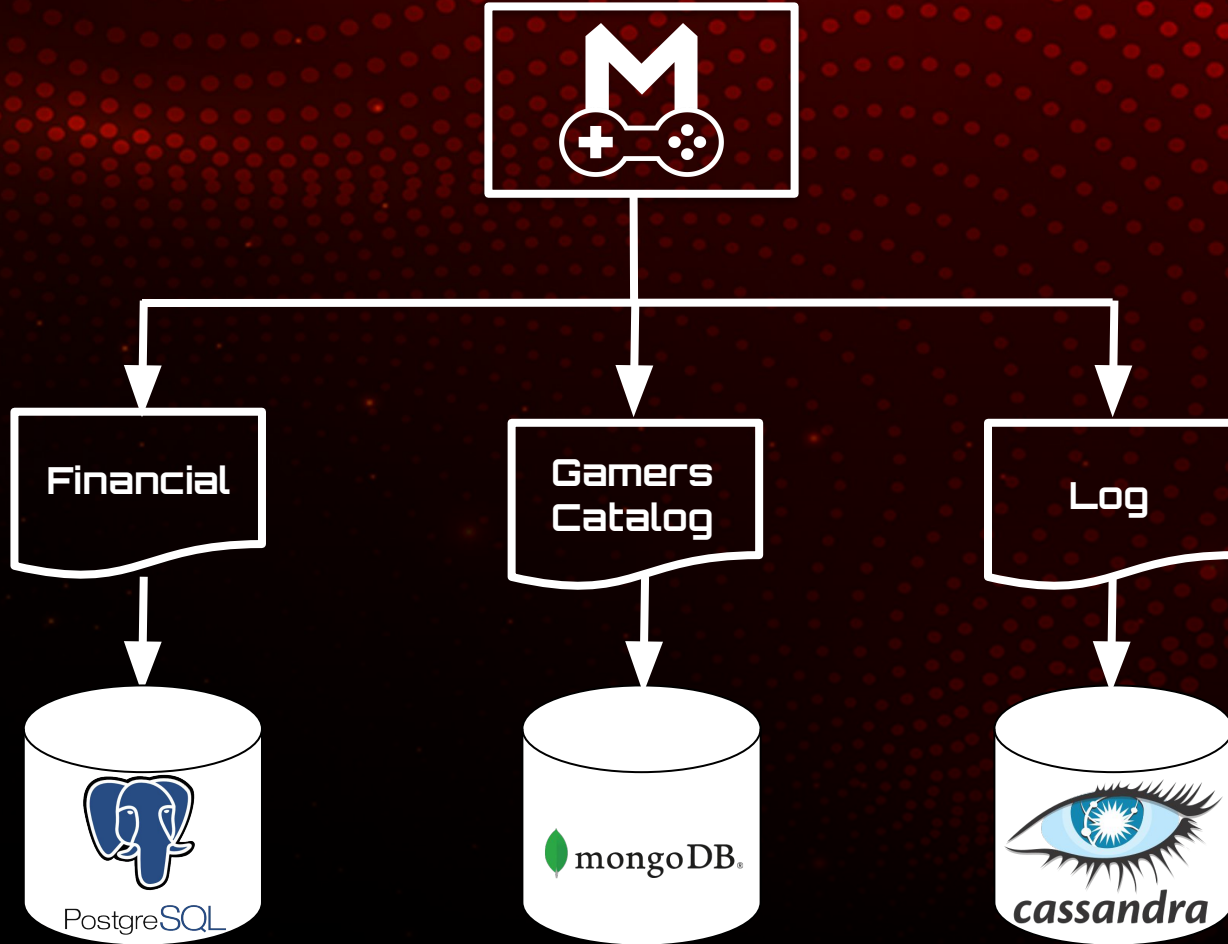
MB Polystore



MB Polystore



MB Polystore



Our Implementation

M8 Client BigDAWG

- Spring Boot



- BigDAWG



- PostgreSQL, SciDB, Accumulo



Our Implementation


Dataset

- **MIMIC II** - *Multiparameter Intelligent Monitoring in Intensive Care II*
- **3 types of data:**
 - **Clinical data** (Relational)
 - **Time-series waveform data** (Array)
 - **Textual medical reports** (Key-Value)

Our Implementation

Execution environment

- Oracle Virtual Box - Ubuntu 18.04
- Docker
 - postgres-catalog
 - postgres-data1
 - postgres-data2
 - scidb
 - accumulo-containers (x4)



Our Implementation

Function: Query

- Executes relational, array or textual queries
- BigDAWG syntax: query token
 - `bdrel(<SQL query content>)`
 - `bdarray(<SciDB query content>)`
 - `bdtext(<Accumulo query content>)`

Demo: Query function

Scegli il tipo di query

- ☒ PostgreSQL
☐ SciDBArray
☐ AccumuloText

Inserisci la query

EXAMPLES:

PostgreSQL: `select * from mimic2v26.d_patients limit 4;`

SciDBArray: `filter(myarray,dim1>150)`

AccumuloText: `{ 'op': 'scan', 'table': 'mimic_logs', 'range': { 'start': ['r_0001',''], 'end': ['r_0015',''] } }`

Esegui Query

Reset

[Back to home](#)

Our Implementation

Function: Catalog

- Interrogate catalog for meta-data knowledge
- BigDAWG syntax: catalog token
 - `bdcatalog(<SQL query content>)`
 - `objects, engines, databases, shims`

Demo: Catalog function

Catalog Objects

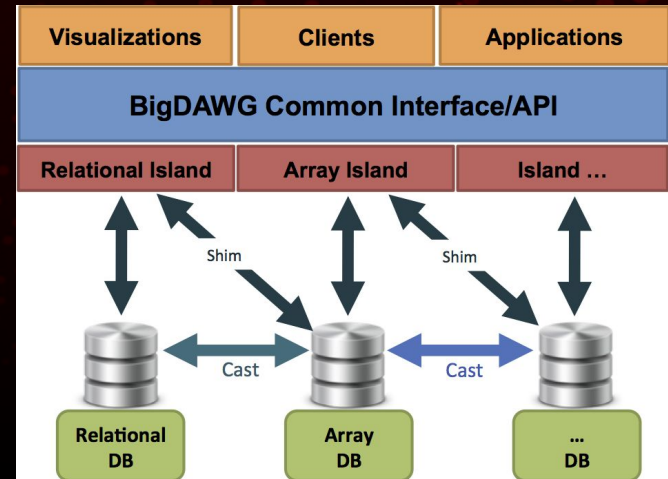
Executed BigDawg Catalog Query: bdcatalog(select e.connection_properties as type, o.logical_db,o.name, o.fields from catalog.objects as o, catalog.engines as e, catalog.databases as d where o.physical_db=d.dbid and d.engine_id=e.eid limit 48)

type	logical_db	name	fields
"PostgreSQL 9.4.5"	2	"mimic2v26.a_iodurations"	"subject_id,icustay_id,itemid,elemid,starttime,startrealtime,endtime,cuid,duration"
"PostgreSQL 9.4.5"	2	"mimic2v26.a_medddurations"	"subject_id,icustay_id,itemid,elemid,starttime,startrealtime,endtime,cuid,duration"
"PostgreSQL 9.4.5"	2	"mimic2v26.additives"	"subject_id,icustay_id,itemid,ioitemid,charttime,elemid,cgid,cuid,amount,doseunits,route"
"PostgreSQL 9.4.5"	2	"mimic2v26.admissions"	"hadm_id,subject_id,admit_dt,disch_dt"
"PostgreSQL 9.4.5"	2	"mimic2v26.censusevents"	"census_id,subject_id,intime,outtime,careunit,destcareunit,dischstatus,los,icustay_id"

Our Implementation

Function: Cast

- Executes cast function
- BigDAWG syntax: annidate cast token `bdcast(...)`
 - Array to Relational
 - Relational to Array
 - Relational to Textual



Demo: Cast function

Scegli il tipo di cast

- ☒ Array->Relational
- ☐ Relational->Array
- ☐ Relational->Text

Inserisci la query di origine

EXAMPLES:

Array->Relational: filter(myarray,dim1>150)

Relational->Array: select poe_id, subject_id FROM mimic2v26.poe_order limit 5

Relational->Text: select * from mimic2v26.icd9 limit 4

Inserisci il nome della tabella di arrivo

Inserisci lo schema della tabella di arrivo

EXAMPLES:

Array->Relational: '(i bigint, dim1 real, dim2 real)'

Relational->Array: '[poe_id=0:*,10000000,0]'

Relational->Text: ''

Esegui Cast

Demo: Query/Cast results

Executed BigDawg Query: bdrel(select * from mimic2v26.d_patients limit 4)

subject_id	sex	dob	dod	hospital_expire_flg
67	"M"	"2903-06-04 00:00:00.0"	"2976-11-29 00:00:00.0"	"Y"
56	"F"	"2553-05-26 00:00:00.0"	"2644-01-23 00:00:00.0"	"Y"
37	"M"	"3195-09-11 00:00:00.0"	"3265-12-31 00:00:00.0"	"N"
78	"M"	"2729-08-08 00:00:00.0"	"2781-03-11 00:00:00.0"	"N"

Executed BigDawg Query: bddarray(scan(bdcast(bdrel(select poe_id, subject_id FROM mimic2v26.poe_order limit 5), prova, '<subject_id:int32>[poe_id=0*,10000000,0]', array)))

poe_id	subject_id
710072	37
710073	37
710075	37
710079	37
710090	37

Future developments

- Integrate MongoDB and Cassandra engines
- Introduce analytic layer for batch processing

Thanks

Thanks for your
attention!

pie.delellis@stud.uniroma3.it | 512521
fra.foresi@stud.uniroma3.it | 508971

